

Research and Planning Report

94-815 Agent Based Modelling and Agentic Technology

Team 8 - Fanxing Bu, Ivan Wiryadi

Project Overview

The base aim for our project is to build an application (portal) using LLM based agents on top of **expense receipt**, called *Sticklet*. The agent is able to extract and store key information from the receipts and use the records to:

1. Visualize trends and data, e.g. weekly or monthly expenses.
2. Discuss with the user about anything relating to their expenses.

1. Overview of Selected Agent Patterns

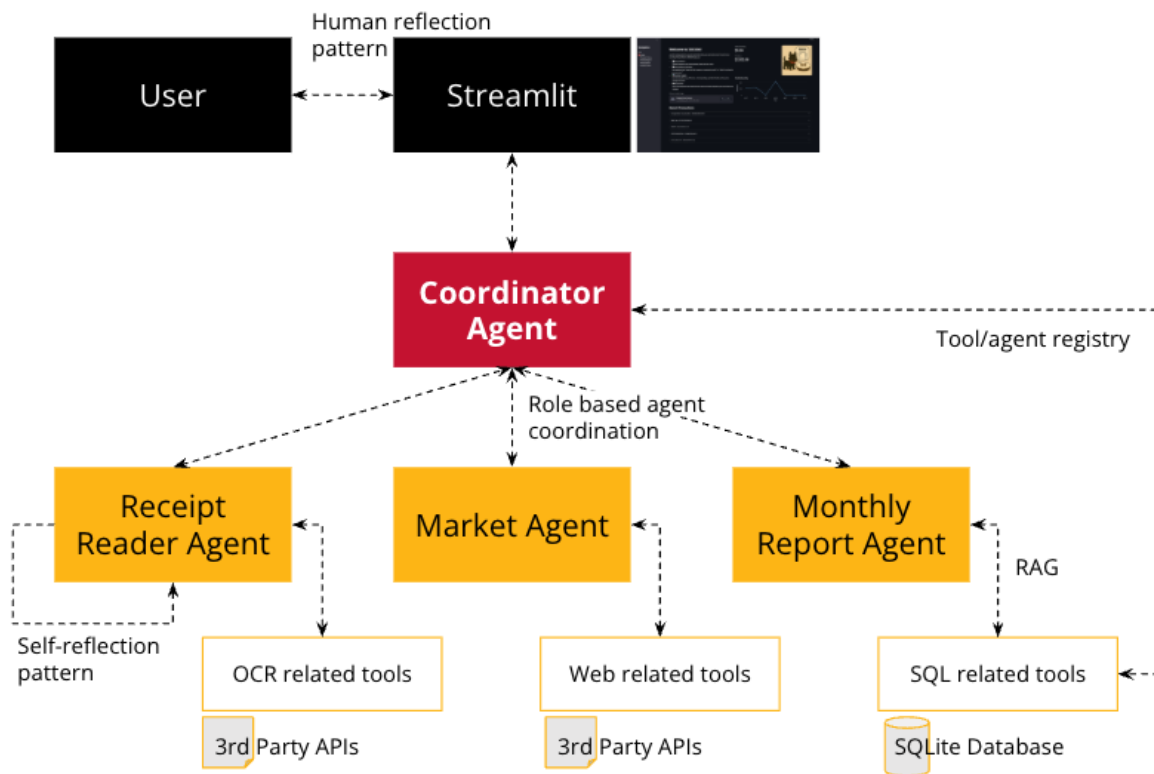


Figure 1: Sticklet Architecture Overview

1.1. Human Reflection Pattern

The Human Reflection pattern in Sticklet enables users to validate and correct AI-processed information through the Streamlit interface. In `app.py`, users can review extracted receipt data (merchant name, date, items, prices), make corrections, and confirm before saving to the database. This pattern increases accuracy by allowing humans to correct mistakes in AI processing, which is particularly important for financial data where errors could lead to incorrect insights.

1.2. Tool/Agent Registry Pattern

The Tool/Agent Registry pattern is implemented in the `CoordinatorAgent` class which maintains a centralized registry of specialized agents and tools. The coordinator uses `_initialize_agents()` and `_get_agent()` methods to lazily initialize and retrieve specialized agents only when needed. Each specialized agent also has access to its required tools - for example, `ReceiptReaderAgent` has OCR and parsing tools, while `MonthlyReportAgent` has access to SQL tools. This pattern enables resource conservation, provides a clean API for agent access, and facilitates modular development.

1.3. Role-Based Agent Coordination Pattern

The Role-Based Agent Coordination pattern distributes specialized tasks to purpose-built agents while the `CoordinatorAgent` orchestrates their activities. Each agent has a distinct role:

- **ReceiptReaderAgent:** Extracts structured data from receipt images
- **MonthlyReportAgent:** Generates financial reports based on spending history
- **MarketAgent:** Retrieves and analyzes market information

The coordinator delegates specific tasks to these specialized agents through methods like `process_receipt()`, `gen_monthly_report()`, and `get_market_indicators()`. This pattern reduces complexity through separation of concerns, allows for specialized optimizations in each agent, and enables parallel development of different agent capabilities.

1.4. Self-Reflection Pattern

The Self-Reflection pattern is implemented in the `ReceiptReaderAgent` through the `_reflect_on_results()` method, which validates and potentially corrects extracted data. When processing receipts, the agent extracts data and then “reflects” on whether the results make sense. For example, it identifies generic merchant names like “receipt” or “store” and attempts to find more specific merchant names from the raw text. It also validates dates and item categories. This pattern improves data quality by identifying and correcting common errors, reducing the need for human intervention, and learning from past mistakes.

1.5. Retrieval Augmented Generation (RAG) Pattern

The RAG pattern enhances the agent’s knowledge with external data from the SQLite database. The system uses `PurchaseMemory` to store structured purchase data and provides tools like `SQLQueryTool` to query this database. When users ask questions about their spending patterns, the agent can retrieve relevant transaction data to provide personalized responses rather than relying solely on its pre-trained knowledge. This pattern enables the foundation model to access domain-specific knowledge, provide personalized responses based on user data, and perform temporal analysis of spending patterns without needing to retrain the model.

2. Tool Comparison and Selection Rationale

To develop Sticklet—a multimodal personal receipt management portal powered by LLM agents—we conducted a detailed evaluation of candidate tools across five key technical components:

- OCR & Information Extraction
- Language Model Reasoning
- Agent Orchestration
- Memory & Persistence
- Financial Data Integration

2.1 OCR and Information Extraction

We compared traditional OCR tools like **Tesseract** and **PaddleOCR** with emerging **Vision-Language Models (VLMs)** like **Mistral**. While conventional pipelines require a separate **Key Information Extraction (KIE)** stage after text recognition, VLMs integrate both steps and support unstructured layouts with better robustness.

We selected **Mistral OCR** due to:

- **Superior multimodal performance:** Robust handling of tilted, cropped, or degraded receipt images.
- **One-shot structured extraction:** Avoids fragile OCR → NLP cascades.
- **API availability:** Currently free, well-documented, and easily replaceable (e.g., GPT-4o or Gemini).

This end-to-end approach reduces pre-processing complexity and improves system resilience.

2.2 Language Model Reasoning

For natural language understanding and reasoning tasks (e.g., agent coordination, summarization), we adopted a hybrid strategy:

- **GPT-4** for high-quality orchestration and instruction-following.
- **Mistral LLM** as a fallback for lightweight or budget-sensitive queries.

This dual-model strategy balances **performance** and **cost-effectiveness**, while maintaining **modularity** and **scalability**. High-stakes reasoning runs on GPT-4; routine extraction or summarization defaults to Mistral.

2.3 Agent Orchestration and Framework

We chose **LangChain** as the primary framework due to:

- **Rich agent tooling** (**AgentExecutor**, **ReAct**, OpenAI function-based agents).
- **Tool integration support** (**BaseTool**, custom tool registration).
- **Prompt management** (**ChatPromptTemplate**, **MessagesPlaceholder**).
- **Memory and context-sharing components** (**SimpleMemory**, **SharedMemory**).

LangChain’s modular, Pythonic design enabled rapid prototyping and extension across the entire multi-agent pipeline.

2.4 Memory and Database

For persistent memory and structured querying, we used:

- **SQLite** for backend data storage.
- **LangChain Memory** tools (**SimpleMemory**, **ReadOnlySharedMemory**) for fast lookup and agent context propagation.

This setup supports seamless integration between reasoning agents and historical data, especially for personalized queries and temporal pattern analysis.

2.5 Financial Market Integration

We adopted a **two-layered approach**:

- **Market data** via the **Yahoo Finance API** (for stock indices, prices).
- **Financial news** using **LLM browsing** or plugin tools (e.g., GPT with browsing tools) for real-time headline summarization.

This balances **structured data retrieval** with **natural language summarization**, keeping the user informed of external trends relevant to their purchases.

2.6 Tool Summary Table

Function Area	Candidates	Final Choice	Rationale
OCR	Tesseract, PaddleOCR, Mistral	Mistral VLM	End-to-end extraction, robust to noise
LLM Reasoning	GPT-4, Mistral	GPT-4 + Mistral	Fidelity + fallback affordability
Agent Framework	LangChain, Autogen	LangChain	Mature ecosystem, prompt/memory/tools
Memory / Persistence	Custom SQLite, LangChain memory	LangChain + SQLite	Rich context + structured data
Market Data & News	Yahoo API, News Scraping, Plugins	Yahoo + LLM	Accurate + low-integration overhead

This toolchain enabled us to quickly build a robust, multimodal MVP with high accuracy, modularity, and extensibility—while staying within a course project’s budget and time constraints.

3. Conceptual Design and Use Cases

The core function of Sticklet is to scan and transform receipt images into structured data that augments the personal financial agent. Users are able to upload receipt images through the web interface, which is then processed to extract information such as merchant, transaction date, total amount, the items purchased, etc and store them into persistent database. Using these data, the agent can help to provide relevant information about the users expenses.

3.1 Use Case: Natural Language Financial Queries

Users can interact with Sticklet through natural language questions about their spending habits and financial patterns. The **CoordinatorAgent** processes these queries by translating them into appropriate database operations using tools like **SQLQueryTool** and **InsightGeneratorTool**. The system can answer questions ranging from simple lookups (“How much did I spend at Trader Joe’s?”) to complex analyses (“How has the price of white rice that I bought changed over time?”). This enables users to gain insights about their finances without needing to know specialized query languages or spreadsheet operations.

3.2 Use Case: Monthly Reporting & Analysis

Sticklet generates comprehensive monthly reports that provide users with an overview of their spending patterns. The **MonthlyReportAgent** aggregates all transaction data for a specified month, calculates totals by category and merchant, identifies high-spend days, and generates narrative summaries using the Mistral API. These reports include both quantitative data visualizations (spending trends, category breakdowns) and qualitative analysis (patterns, anomalies, recommendations), giving users actionable insights about their financial behavior without requiring manual data compilation.

3.3 Use Case: Market Intelligence Integration

The Market Agent connects personal finance data with broader market context by tracking major market indices (S&P 500, Dow Jones, NASDAQ) and generating tailored market summaries. Users can view 7-day historical market data and receive AI-generated narratives that explain market movements. This integration helps users understand how external economic factors might impact their personal finances and make more informed decisions about future spending or investments.

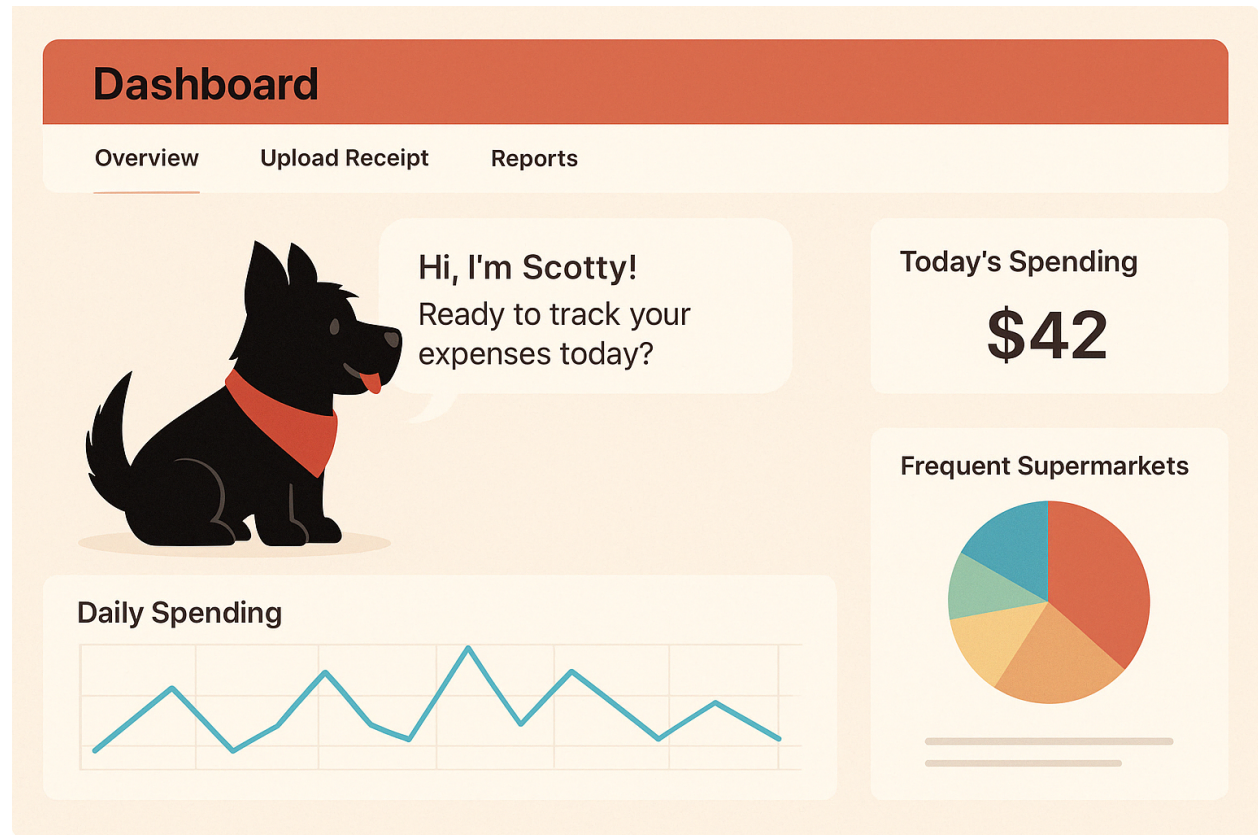


Figure 2: Early design